

Active vision for dexterous grasping of novel objects

Ermano Arruda

Marek Kopicki

Jeremy L. Wyatt

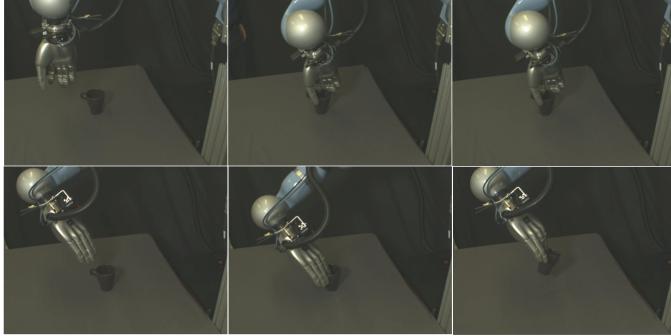


Fig. 1. Top three pictures show a failed grasp due to unexpected collision with parts of the object that are not involved in the grasp. Bottom pictures demonstrate a successful safe grasp trajectory selected after running our proposed approach, that takes into account both search for good contact quality and safety of grasp execution.

Abstract—How should a robot direct active vision so as to ensure reliable grasping? We answer this question for the case of dexterous grasping of unfamiliar objects. When an object is unfamiliar, much of its shape is by definition unknown. An initial view will recover only some surfaces, leaving most of the object’s surface unmodelled, and also leaving shadow regions which may or may not contain obstacles. These two features make it difficult both to select reliable grasps, and to plan safe reach-to-grasp trajectories. Grasps typically fail in one of two ways, either unmodelled objects in the scene cause collisions, or object reconstruction is insufficient to ensure that the grasp points provide a stable force closure. These problems can be solved more easily if active sensing is guided by the anticipated actions. Our approach has three stages. First, we take a single view and generate candidate grasps from the resulting partial object reconstruction. Second, we drive active vision to maximise surface reconstruction quality around the planned contact points. During this phase the anticipated grasp is continually refined. Third, we direct gaze to unmodelled regions that will affect the planned reach to grasp trajectory, so as to confirm that this trajectory is safe. We show, on a dexterous manipulator with camera on wrist, that our approach (85.7% success rate) outperforms a randomised algorithm (48% success rate).

I. INTRODUCTION

Grasping of novel objects is a hard problem on which there has been steady progress [10], [11], [8], [14], [7], [16], [6], [3], [15], [4]. We now possess methods that are able to generate dexterous grasps for unfamiliar objects, using

We gratefully acknowledge support of FP7 grant IST-600918, PacMan, and a studentship from Brazilian Science without Borders for Ermano Arruda.

Arruda, Kopicki and Wyatt at CN-CR, University of Birmingham, Edgbaston, Birmingham, United Kingdom, B15 2TT, Tel.: +44-121-4144788, Fax: +44-121-4144281, exa371.m.s.kopicki, jlw@cs.bham.ac.uk

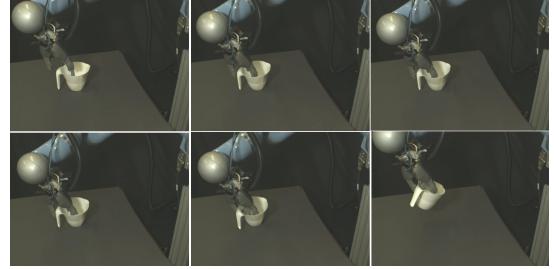


Fig. 2. Sample grasp result from random-based approach. The grasp is successful. However, it can be seen that the grasp trajectory starts moving and pushing the object aside with its fingers long before the final grasp closure takes place in the last picture on the right. This is a typical scenario which could lead to failed grasp, due to additional imprecision, different shaped objects, or damage to the hand.

incomplete object reconstructions. Nonetheless the reliability of grasping rises with the quality and completeness of the reconstruction available. Given an active vision system, we would like to minimise the number of views taken, while maximising grasping reliability.

At the root of the difficulties is a chicken and egg problem. On the one hand, given that the initial point cloud can be highly incomplete, it is hard to plan a reliable grasp to begin with. On the other hand if we knew the likely planned grasp then we could direct gaze more efficiently. In this paper we solve this problem by employing a grasp planner that can generate grasps for novel objects in the face of fragmentary reconstructions. We use grasp candidates to guide active vision, and the results of active vision to refine grasp planning.

First we describe related work, and then proceed to describe our active vision method. This has two parts, a routine driven by the planned contact points, and a routine driven by the need to ensure a safe reach to grasp trajectory. We then present experimental results on 14 novel objects, comparing our method with a randomised view planner.

II. RELATED WORK

Active vision, or more generally active perception, is defined as the study of modelling, planning and control strategies for perception when the sensor can be actively moved [2]. The field of active perception for robots started with work by [2], [1]. The greatest advantage of active perception is that many problems that are hard to solve in the passive observer paradigm become easier. For example, classical computer vision problems such as *structure from motion* and *optical flow* were shown to be well behaved if an agent integrates visual information over time while

performing controlled motion [1]. Since then, in the context of manipulation, researchers have focused on devising strategies for view selection based on recovery of the full shape of the object to be grasped [12], [5]. For most practical manipulation purposes, however, full object reconstruction is too costly or simply infeasible. It is also typically redundant, since most of the time only a limited portion of the object surface is in contact during a grasp [10].

III. VIEW SELECTION

We now sketch our method, before proceeding to the details. The robot begins by taking a single view from a fixed location of the scene. A depth camera is used, and this gives an incomplete point cloud. A dexterous grasp planning algorithm is then run, which generates a large number of candidate grasps on the partial point cloud for the object. These grasps will typically assume the existence of graspable surfaces on both sides of the surface defined by the point cloud. The predicted contact locations are then used to drive the next view. The next view is chosen to maximise the quality of the point cloud at the planned contact locations. If a grasp cannot be found, we employ information gain view planning, using 3D occupancy map. Once the quality of the relevant surface reconstruction is sufficiently high, or a limit on the number of views is reached, the grasp is fixed. Then a second phase of active vision aims to verify a safe path to the grasp location. To achieve this we again use the 3D occupancy map. This is used to calculate the probability of a collision free trajectory. Active views for safety are driven to reduce the average entropy in cells through which the candidate reach-to-grasp trajectory passes. This ensures a safe grasp. We now proceed to describe the representations, and the three criteria used to drive active vision at different stages (contact based, information gain, and safety based).

A. Representations

We start by describing the underlying representations used to define our approach. Let $\Xi = [\xi_1, \xi_2, \dots, \xi_N]$ be a list of hypothesis of camera poses, where $\xi_i \in SE(3)$, and $V \subset \Xi$ a set of already visited camera poses. In addition, let γ be a point cloud obtained from a certain camera pose ξ . We define Γ_t as our object point cloud segmented out of the table plane, after a number of t views have been taken,

$$\Gamma_t = \Gamma_{t-1} \uplus \text{segmented}(\gamma), \quad (1)$$

i.e., Γ_t is the result of segmenting the object point cloud from the table plane in γ and integrating this result with our previous obtained object point Γ_{t-1} .

In addition to the object cloud, we also maintain a more complete representation of the full robot workspace as a 3D occupancy grid implementation with OcTrees by [9]. We shall refer to this representation of the robot workspace as Λ , which is updated at every measurement (ξ, γ) . This particular implementation allow us to easily represent known and unknown parts of the robot workspace Λ and to define information gain view selection strategies, by reasoning on

top of the occupancy probabilities in the subsequent sections as we shall see.

It is possible to find a grasp trajectory τ by transferring a learnt grasp G to the given object represented by Γ_t using the method of Kopicki, Wyatt et al. [10]. And using the same method, we are able to extract contact points from Γ_t , yielding a list of contacts $C = [\mathbf{c}_1, \dots, \mathbf{c}_M]$, where $\mathbf{c}_i = (w_i, \mathbf{p}_i, \mathbf{n}_i)$ is composed of a weight $w_i \in \mathcal{R}$, indicating its relevance to the grasp, the contact location $\mathbf{p}_i \in \mathcal{R}^3$, and the surface normal at that point $\mathbf{n}_i \in \mathcal{R}^3$.

Let us also define $\Omega = [(\mathbf{c}_1, z_1), \dots, (\mathbf{c}_M, z_M)]$ as a list of contact points expanded to include the current quality of the observation of that point from the best view ξ to date. We proceed now to describe a contact-based view selection approach.

B. Contact Based View Selection

If we let the viewing direction of a certain view pose $\xi_k \in \Xi$ be the vector \mathbf{v}_k always pointing towards the object Γ_t , then we can define the quality of observation of a contact point \mathbf{c}_i from a given ξ_k as

$$\theta_{ki} = \theta(\xi_k, c_i) = \arccos(\min(0, \mathbf{v}_k^T \mathbf{n}_i)). \quad (2)$$

Observe that when looking at a certain contact point with a given normal \mathbf{n}_i , we assign higher values to views that look at certain contact point directly, i.e. \mathbf{n}_i and \mathbf{v}_k form an angle of 180 degrees, according to our convention that \mathbf{v}_k always looks towards the object.

Thus, for each element $(\mathbf{c}_i, z_i) \in \Omega$ we store the contact points \mathbf{c}_i , and define z_i the best quality of observation to date with respect to all visited poses as

$$z_i = \arg \max_{\xi_j \in V} \theta_{ji}. \quad (3)$$

Note that Eq 2 and Eq 3 express our desire to prefer views in which the observation of the object surface is maximally reliable with respect to depth error. The intuition behind this heuristic is that, depth cameras typically yield poor point cloud data when the surface is observed obliquely with respect to the surface normal.

Finally, let $F_\tau = [\mathbf{f}_1, \dots, \mathbf{f}_R]$ be the list of finger link normals at the last time step of the trajectory τ . We then define the individual value of choosing a given non-visited view to look at a particular contact \mathbf{c}_i as

$$\sigma(\xi_k, F_\tau, \mathbf{c}_i) = w_i \sum_{r=1}^R \max(\theta_{ki}, z_i) \frac{(1 - \text{sign}(\mathbf{f}_r^T \mathbf{n}_i))}{2}. \quad (4)$$

Observe that in Eq 4, when looking at a certain contact point \mathbf{c}_i we are able either to improve our previous best viewing quality if $\theta_{ki} > z_i$, or leave it as it is. Note also that the multiplying term $\frac{(1 - \text{sign}(\mathbf{f}_r^T \mathbf{n}_i))}{2}$ serves as a switch that turns this value relevant or not, yielding 0 or 1, according to the geometric constraint that a link must have with a contact point, i.e. the surface normal \mathbf{f}_r of a given finger link must point in the opposite direction of the contact normal \mathbf{n}_i , otherwise this contact point is meaningless with respect to

this given finger link, meaning that viewing it is not useful. Finally, the normalised weight w_i scales this value according to its overall relevance to the grasp as defined by Kopicki, Wyatt et al. [10] approach. It follows that the total utility value of a given view ξ is given by

$$u_1(\xi, \Omega, \tau) = \sum_{i=1}^N \sigma(\xi, F_\tau, \mathbf{c}_i). \quad (5)$$

We are then able to rank next potential views by calculating the total value of a view with respect to all contact points, and picking the view that has maximum such value as in Eq 6.

$$\xi^* = \arg \max_{\xi_k \in \Xi - \mathbf{V}} u_1(\xi, \Omega, \tau). \quad (6)$$

C. Information Gain View Selection

Unfortunately, when no contact points are found in a given view, the previous view selection procedure will fail. When such case happens, we define an information-gain based utility function for view selection. Intuitively, this strategy makes sense, since no contacts were found with the knowledge we have about the object shape so far, represented by Γ_t . Therefore, one should ideally adopt a completely exploratory behaviour to seek for new parts of the object.

For this purpose, let $bmin(\Gamma_t), bmax(\Gamma_t) \in \mathcal{R}^3$ be the respective minimum and maximum limits of the bounding box that circumscribes the object point cloud Γ_t . We are able then to extract the set of voxels $\Lambda_{object} \subset \Lambda$ inside this bounding box. If we assume the surface of this voxelised solid box Λ_{object} is visible from all cameras, as shown in Fig 4. Then we can define a simple strategy to minimise the entropy about the object's shape, by selecting views that are going to have maximum predicted information gain about the voxels in Λ_{object} . Intuitively, our goal is to select views such that we gradually sculpt the solid cube, such that we will eventually reach a constant entropy value for this cube, due to self-occluding parts of the object, from which point no views are going to bring any more information gain.

Our first step to fulfil this goal is to define a rule with which we can tell the set of visible voxels in Λ_{object} visible from a camera pose ξ . Such visibility test is performed as an implementation of the typical frustum culling graphics procedure [InsertRef], with a few slight modifications. Firstly, as customarily, we transform the set of voxels Λ_{object} into the camera coordinate system. Finally, during the projection phase of the pipeline, we allow free voxels to be projected onto the same image coordinates, but we do not allow either an unknown or an occupied voxels to be projected on top of the other. As a consequence, we end up finding a border in our initial solid cube $\Lambda_{visible}(\xi) \subset \Lambda_{object}$, which contains all free projected voxels onto the image plane, together with the boundary voxels that might be either unknown or occupied, as shown in Fig 3. Thus, $\Lambda_{visible}(\xi) = \{s_1, \dots, s_D\}$ is defined as our set of voxels of interest for information gain prediction. We then follow to describe the information gain prediction for the set of voxels $\Lambda_{visible}(\xi)$.

1) Information Gain Prediction: Let the occupancy probability of occupancy of a voxel $s_d \in \Lambda_{visible}$ up to our most recent observation $o_{1:t-1}$ for this particular be $p_{s_d} = p(s_d|o_{1:t-1})$. We can write the entropy of this voxels by viewing it as a Bernoulli random variable as

$$H(s_d) = -p_{s_d} \log(p_{s_d}) - (1 - p_{s_d}) \log(1 - p_{s_d}), \quad (7)$$

By using a log-odds representation of our occupancy probability such as in [9], [13], we can then define future predicted measurement on s_d as

$$L(s_d|o_{1:t-1}, o'_t) = L(s_d|o_{1:t-1}) + L(s_d|o'_t), \quad (8)$$

where $o'_t \in O = \{\text{occupied}, \text{free}\}$ is an imaginary measurement and $L(s_d|o'_t)$ is also called *inverse sensor model* [17]. The inverse sensor model is defined likewise as in [9] as

$$L(s_d|o) = \begin{cases} L_{occ}, & \text{if } o = \text{occupied}. \\ L_{miss}, & \text{otherwise.} \end{cases} \quad (9)$$

Note that our imaginary occupancy probability converted from log-odds is then $p_{sd|o'_t} = p(s_d|o_{1:t-1}, o'_t) = \frac{1}{1 + \exp(L(s_d|o_{1:t-1}, o'_t))}$.

We make a simplifying assumption that an imaginary measurement has uniform distribution, i.e. $p(\text{occupied}) = p(\text{free}) = 0.5$. Thus, we define our predicted entropy resulting from an imaginary measurement as the expected value

$$H'(s_d) = - \sum_{o' \in O} p(o') (p_{sd|o'} \log(p_{sd|o'}) + (1 - p_{sd|o'}) \log(1 - p_{sd|o'})) \quad (10)$$

Therefore, the information gain of looking at a particular voxel $s_d \in \Lambda_{visible}(\xi)$ from a given view ξ is given by

$$I(\xi, s_d) = H(s_d) - H'(s_d), \quad (11)$$

where the total average information gain per voxel is given by

$$u_2(\xi, \Lambda_{visible}(\xi), \Gamma_t) = \sum_{s_d \in \Lambda_{visible}} \frac{I(\xi, s_d)}{D}, \quad (12)$$

where $D = |\Lambda_{visible}|$. Note that we refer to the average information gain per voxel, since different views have different numbers of visible voxels in their field of view after frustum culling. This makes the predicted information different gain for different views comparable.

2) Information Gain View Selection: Using the definitions aforementioned, when no contacts are available, we are finally able to select next best views according to a maximum information gain strategy via

$$\xi^* = \arg \max_{\xi_k \in \Xi - \mathbf{V}} u_2(\xi_k, \Lambda_{visible}(\xi), \Gamma_t). \quad (13)$$

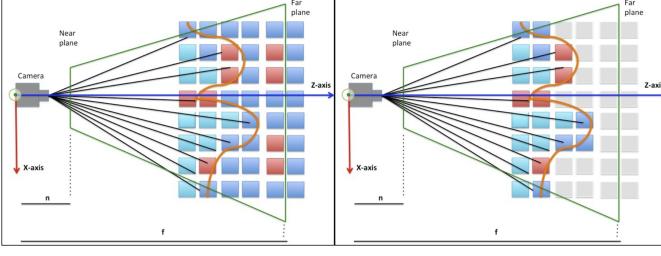


Fig. 3. Cross-section view of a typical visibility check. In the picture, occupied voxels are represented in red, free voxels are cyan and unknown voxels have dark-blue colour. Having defined a viewing frustum to match the real depth camera specifications, frustum culling procedure is performed in which free voxels are assumed to be transparent, whereas unknown or occupied voxels occlude each other. As a result, only the voxels colored on the right image are defined as being visible after the execution of this procedure.

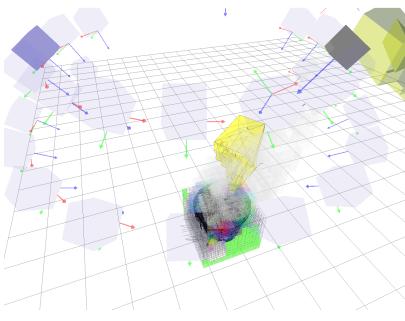


Fig. 4. View camera poses forming the set Ξ , camera pose highlighted in darker purple belongs to the set of visited poses V . Object in the center is circumscribed by a voxelised cube. View exploration sculpts this cube when no contacts are found according to maximum information gain strategy.

Algorithm 1 Grasp Driven Active Sense

```

1: procedure ACTIVEGRASP( $\Xi, G$ )
2:    $\Gamma = \emptyset, \Lambda = \emptyset$ 
3:    $V = \emptyset, T = \emptyset$      $\triangleright$  List of visited views and found
                           trajectories, respectively. Initially empty
4:    $grasp = false$ 
5:    $\tau^* = None$ 
6:   while not  $grasp$  do
7:      $V, \tau^*, \Gamma, \Lambda = nextBestView(\Xi, \Gamma, \Lambda, G, V, T)$ 
8:      $V, p, \Gamma, \Lambda = safetyExploration(\Xi, \Gamma, \Lambda, \tau^*, V)$ 
9:     if  $p \leq \alpha$  then
10:       $grasp = true$ 
11:    else
12:       $T = T - \{\tau^*\}$      $\triangleright$  Removing  $\tau^*$  from our
                           candidate trajectories
13:    end if
14:   end while
15:    $executeGrasp(\tau^*)$ 
16: end procedure
```

Algorithm 2 Next Best View Exploration

```

1: function NEXTBESTVIEW( $\Xi, \Gamma, \Lambda, G, V, T$ )
2:    $\Omega = \emptyset$            $\triangleright$  Most recent found contact points
3:    $\tau = None$           $\triangleright$  Most recent found grasp trajectory
4:    $stop = false$ 
5:   while not  $stop$  do
6:      $\xi = selectNBV(\Xi, \Gamma, \Lambda, V, \Omega, \tau)$ 
7:      $V = append(V, \xi)$      $\triangleright$  Appending  $\xi$  to  $V$ .
8:      $\gamma = capture(\xi)$      $\triangleright$  Point cloud from pose  $\xi$ .
9:      $\Gamma = \Gamma \uplus segmented(\gamma)$ 
10:     $\tau, \Omega = findGrasp(\Gamma, G)$ 
11:     $\Lambda = updateOctree(\Lambda, \gamma, \xi, \Omega)$ 
12:     $T = append(T, \tau)$ 
13:     $stop = CHECKSTOP(V, T)$ 
14:   end while
15:    $\tau^* = \arg \min_{\tau \in T} p(\tau | \Lambda)$ 
16:   Return  $(V, \tau^*, \Gamma, \Lambda, T)$ 
17: end function
18: function CHECKSTOP( $V, T$ )
19:   Return  $(|V| \geq 2 \text{ and } |T| \geq 1) \text{ or } |V| \geq 7$ 
20: end function
```

Algorithm 3 Select Next Best View Contact Based

```

1: function NBVCONTACTBASED( $\Xi, \Gamma, \Lambda, V, \Omega, \tau$ )
2:    $\xi = None$ 
3:   if  $|V| = 0$  then
4:      $\xi = head(\Xi)$ 
5:   else if  $\Omega \neq \emptyset$  then
6:     Let  $\xi$  be selected according to Eq 6
7:   else
8:     Let  $\xi$  be selected according to Eq 13
9:   end if
10:  Return  $\xi$ 
11: end function
```

Algorithm 4 Select View Randomly

```

1: function SELECTNBVRANDOM( $\Xi$ )
2:    $\xi = None$ 
3:   if  $|V| = 0$  then
4:      $\xi = head(\Xi)$      $\triangleright$  First element of the list  $\Xi$ 
5:   else
6:      $\xi = randomUniformSample(\Xi)$ 
7:   end if
8:   Return  $\xi$ 
9: end function
```

Algorithm 5 Safety Exploration

```

1: function SAFETYEXPLORATION( $\Xi, \Gamma, \Lambda, \tau, V$ )
2:    $stop = false$ 
3:   while not  $stop$  do
4:      $\xi, value = safetyNBV(\Xi, \Lambda, \tau)$ 
5:      $V = append(T, \xi)$ 
6:      $\gamma = capture(\xi)$        $\triangleright$  Point cloud from pose  $\xi$ 
7:      $\Gamma = \Gamma \uplus segmented(\gamma)$ 
8:      $\Lambda = updateOcTree(\Lambda, \gamma, \xi, \Omega)$ 
9:      $T = append(T, \tau)$ 
10:     $stop = CHECKSTOPSAFETY(value)$ 
11:   end while
12:    $p = p(\tau | \Lambda)$ 
13:   Return  $(V, p, \Gamma, \Lambda)$ 
14: end function
15: function CHECKSTOPSAFETY( $value$ )
16:   if  $value \leq \beta$  then
17:     Return true
18:   else
19:     Return false
20:   end if
21: end function

```

Algorithm 6 Safety Exploration View Selection

```

1: function SAFETYNBV( $\Xi, \Lambda, \tau$ )
2:    $\Lambda_c = findCollisionVoxels(\Lambda, \tau)$ 
3:   Using  $\Lambda_c$ , let  $\xi$  be selected according to Eq 13
4:    $value = u_2(\xi, \Lambda_c)$ 
5:   Return  $(\xi, value)$ 
6: end function

```

D. Safety View Planning

In the safety view planning we are interested in estimating the probability of collision during the execution of a given trajectory τ , disregarding the collision with the final contact points Ω . Effectively, we estimate the probability of an unexpected collision along the trajectory τ . This is a typical scenario in which the robot hand collides with an unknown part of the object due to the fact that the grasp was originally planned from an incomplete model of the object's shape Γ_t . In addition, we are also able to access how certain we are regarding this collision estimation by computing the current entropy for this particular trajectory τ . As such, we select views as to minimise the entropy of the voxels through which the robot hand is going to pass when following a given grasp trajectory τ . This enable us to have a final relatively certain estimation with respect to unexpected collisions that might damage the robot hand, or simply make the grasp fail.

Let the set of voxels through which the hand bounds passes when following a trajectory τ be Λ_c . These voxels are retrieved by simulating the hand moving along the trajectory τ and querying at each time step of this trajectory the voxels with which the hand is colliding in our voxelised workspace

Λ . Having retrieved those collision voxels, let p_{s_c} be the probability of occupancy of a given voxel $s_c \in \Lambda_c$. The probability of collision can be calculated by

$$p_{collision}(\tau, \Lambda_c) = 1 - \prod_{s_c \in \Lambda_c} (1 - p_{s_c}). \quad (14)$$

For numerical reasons, we prefer to refer to Eq 14 using only the product term, representing the probability that all voxels along the trajectory τ are free, in its logarithmic form as

$$\kappa(\tau, \Lambda_c) = \ln \prod_{s_c \in \Lambda_c} (1 - p_{s_c}) = \sum_{s_c \in \Lambda_c} (1 - p_{s_c}). \quad (15)$$

Note that $p_{collision}(\tau, \Lambda_c) = 1 - \exp(-\kappa(\tau, \Lambda_c))$.

Finally, to select views in order to get better estimations for Eq 14, we use the same utility function defined in 12. Thus if we let $\Lambda_{c_{visible}}(\xi) \subset \Lambda_c$ be the set of visible voxels by a certain view pose ξ . Next best views are then selected according to

$$\xi^* = \arg \max_{\xi_k \in \Xi - V} u_2(\xi_k, \Lambda_{c_{visible}}(\xi), \Gamma_t). \quad (16)$$

In practise, we allow safety exploration to run while the information gain is above a predefined threshold, i.e. $u_2(\xi_k, \Lambda_{c_{visible}}(\xi), \Gamma_t) > \beta$. If this criteria is not met, the final probability of collision is reported according to Eq 14. The trajectory τ is therefore executed or not based on the probability of collision.

IV. EXPERIMENTS

The following section outlines the experiments we conducted to test our view selection approach. A general pseudo-code of our implementation is described in Alg 1, which is divided into two sub-phases. First a contact-based next best view exploration procedure is run as outlined by Alg 2. In this first phase, at least two views are selected, up to a maximum of 7 views if after the second view no grasp trajectory and contacts are found. The first view is also fixed, only subsequent views after this fixed view are selected according to the criteria described in our previous sections for contact-based view selection. The second phase of Alg 1 is run in order to estimate the probability of collision of the best promising candidate trajectory τ^* , selected as the trajectory with lowest probability of collision prior to the safety view exploration phase, given our current knowledge of the object Γ_t and workspace Λ . This second phase is outlined by Alg 5. Note that the safety exploration phase stops if the current selected view has predicted information gain below a certain threshold β . If at the end of the safety exploration phase, we discover that this trajectory τ^* has probability of collision above a certain threshold α , we reject this trajectory and cycle back to phase 1, i.e. Alg 2.



Fig. 5. Picture shows the 14 objects used for trials.

A. Methodology

Using Alg 1 we performed trials on a set of 14 objects shown in Fig 5. In our experiments, we performed a comparison of our full approach with a modified approach whose sole criteria for view selection in any phase is a random selection strategy. In other words, we substituted all calls of the selection procedures Alg 3 and Alg 6 by the random view selection procedure outlined in Alg 4. Furthermore we limited the two phases of this modified random-based approach to be constrained to the same number of views that our algorithm performed in both phases. It is also important to note that in our experiments we have set the size of the voxels in our 3D occupancy map Λ to be $0.0025m$, for a relatively fine precision. The data for this experiment was collected and can be seen in Table I.

B. Results

The results shown in Table I outline the contrast between the two approaches. We first note that the success rate of our proposed view selection approach achieved a success rate of 85.7%, whereas the modified random-based approach showed a success rate of only 48%. A closer look at Table I reveals that random exploration tended to yield unsafe grasps, under the same view number constraints as our original approach. This indicates that random view selection would probably need to cycle back to generate new grasp trajectory candidates more times, which seems a natural consequence of its sub-optimum exploratory behaviour. One such example is highlighted by Fig 1, in which the final trajectory executed with probability of collision 1.0 and, indeed, makes the robot hand collide with a part of the mug not involved in the grasp, finally failing for safety reasons. We also note that our collision probability appears to be over-sensitive, the random approach also succeeded for various cases in which the probability of collision was 1.0. Nonetheless, even for successful grasps as the one depicted in Fig 2, grasps with probability 1.0 tended to collide prematurely with different parts of the object. In addition, we also noted that for the case of the toothpaste, the trivial solution of a grasp with as few collisions as possible might yield grasps with very poor grip. This indicates future work towards a middle ground between this two extremes of a grasping trajectory.

Finally, as shown by Table I and in Fig 6, our approach had equivalent success rate to prior work done by Kopicki

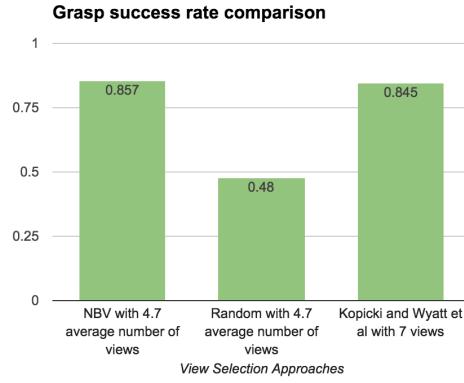


Fig. 6. Grasp success rate comparison.

and Wyatt et al [10]. In our experiments our approach used in average 4.7 views for grasp planning. Even though subsequent views were taken to estimate the collision probability of a given grasp, the grasp itself was planned on a lower number of views, prior to the safety exploration phase, than the fixed number of 7 views performed by Kopicki and Wyatt et al.

TABLE I
TRIAL RESULTS

Objects	Phase 1 View Count	Phase 2 View Count	Active Grasp Result	Collision Probability	Grasp Planning Views
NBV and Random	NBV and Random	NBV and Random	Random	NBV	Random
bowl	4	4	success	fail	0.00009
bowl small	3	4	success	success	0.00104
book	5	4	success	success	0.000035
coke	2	5	success	success	0.002384
cup1	3	5	success	success	0.001015
dustpan	3	3	success	success	0.002267
glass2	4	4	fail	fail	1.0
guttering	3	4	success	success	0.001526
leg	3	5	success	success	0.000057
mmuscle	3	4	success	success	0.003768
mug1	3	4	success	fail	0.0009
remie	3	3	success	success	0.00313
stand2	3	3	success	success	0.006257
toothpaste	4	4	fail	fail	0.00491
			Success Rate	0.857	Average Number of Views
				0.48	4.7

V. CONCLUSIONS

We have proposed an effective approach for view selection composed of two stages. The first stage primarily exploits contact points, seeking for low noise point cloud measurement in an heuristic fashion. If no contact points are available for such exploration to take place, we adopt a purely exploratory behaviour guided by the minimisation of the entropy related to a given object defined in our 3D occupancy grid map as a solid cube to be sculpt around our object cloud Γ_t . After a certain number of candidate trajectory is found, our approach is able to investigate the safety of a good grasp trajectory candidate, accepting or rejecting once more information about this trajectory surroundings is found. We showed that our approach yields better success rate performance when compared to a random-based strategy. We also showed that we slightly improve previous performance results that employed a fixed number of 7 views by Kopicki and Wyatt et al [10]. As future work, we intend to explore the collision sensitivity capabilities of the system, towards an autonomous and active grasping agent.

REFERENCES

- [1] John Aloimonos, Isaac Weiss, and Amit Bandyopadhyay. Active vision. *International Journal of Computer Vision*, 1(4):333–356, 1988.
- [2] Ruzena Bajcsy. Active perception. In *Proc IEEE*, 76:996–1005, 1988.
- [3] H. Ben Amor, O. Kroemer, U. Hillenbrand, G. Neumann, and J. Peters. Generalization of human grasping for multi-fingered robot hands. In *IEEE/RSJ International Conference on Intelligent Robots and Systems*, 2012.
- [4] Jeannette Bohg, Matthew Johnson-Roberson, Beatriz Len, Javier Felip, Xavi Gratal, N Bergstrom, Danica Kragic, and Antonio Morales. Mind the gap-robotic grasping under incomplete observation. In *IEEE International Conference on Robotics and Automation*, pages 686–693. IEEE, 2011.
- [5] Shengyong Chen, Youfu Li, and Ngai Ming Kwok. Active vision in robotic systems: A survey of recent developments. *I. J. Robotic Res.*, 30:1343–1377, 2011.
- [6] Noel Curtis and Jing Xiao. Efficient and effective grasping of novel objects through learning and adapting a knowledge base. In *IEEE/RSJ International Conference on Intelligent Robots and Systems*, pages 2252–2257. IEEE, 2008.
- [7] Renaud Detry and Justus Piater. Unsupervised learning of predictive parts for cross-object grasp transfer. In *IEEE/RSJ International Conference on Intelligent Robots and Systems*, 2013.
- [8] Martin Hjelm, Renaud Detry, Carl Henrik Ek, and Danica Kragic. Representations for cross-task, cross-object grasp transfer. In *IEEE International Conference on Robotics and Automation*, 2014.
- [9] Armin Hornung, Kai M. Wurm, Maren Bennewitz, Cyrill Stachniss, and Wolfram Burgard. Octomap: an efficient probabilistic 3d mapping framework based on octrees. *Autonomous Robots*, 34(3):189–206, 2013.
- [10] Marek Kopicki, Renaud Detry, Maxime Adjigble, Rustam Stolkin, Ales Leonardis, and Jeremy L. Wyatt. One-shot learning and generation of dexterous grasps for novel objects. *The International Journal of Robotics Research*, 2015. first published on September 18, 2015.
- [11] Marek Kopicki, Renaud Detry, Florian Schmidt, Christoph Borst, Rustam Stolkin, and Jeremy L. Wyatt. Learning dextrous grasps that generalise to novel objects by combining hand and contact models. In *IEEE International Conference on Robotics and Automation*, pages 5358–5365. IEEE, 2014.
- [12] Michael Krainin, Brian Curless, and Dieter Fox. Autonomous generation of complete 3d object models using next best view manipulation planning. In *Robotics and Automation (ICRA), 2011 IEEE International Conference on*, pages 5031–5037. IEEE, 2011.
- [13] H. Moravec and A. Elfes. High resolution maps from wide angle sonar. In *Robotics and Automation. Proceedings. 1985 IEEE International Conference on*, volume 2, pages 116–121, Mar 1985.
- [14] Alexander Rietzler, Renaud Detry, Marek Kopicki, Jeremy L. Wyatt, and Justus Piater. Inertially-safe grasping of novel objects. In *Cognitive Robotics Systems: Replicating Human Actions and Activities (Workshop at IROS 2013)*, 2013.
- [15] A. Saxena, L. Wong, and A.Y. Ng. Learning grasp strategies with partial shape information. In *Proceedings of AAAI*, pages 1491–1494. AAAI, 2008.
- [16] Michael Stark, Philipp Lies, Michael Zillich, Jeremy Wyatt, and Bernt Schiele. Functional object class detection based on learned affordance cues. In *Computer Vision Systems*, pages 435–444. Springer, 2008.
- [17] Sebastian Thrun, Wolfram Burgard, and Dieter Fox. *Probabilistic Robotics (Intelligent Robotics and Autonomous Agents)*. The MIT Press, 2005.